



GPU Programming in **julia**

Neil Gutkin

Brief History




Alan Edelman



Jeff Bezanson

- Work started in 2009, first release in 2012
 - Developed and incubated at MIT
 - Jeff Bezanson, Stefan Karpinski, Viral B. Shah, and Alan Edelman
 - Open-source and free under MIT license
 - Julia 1.7.0 was released November 30, 2021



Stefan Karpinski



Viral B. Shah


- Why?
 - “In short, because we are greedy” [1]
 - Wanted a fast, high-level, all-purpose language



[1] <https://julialang.org/blog/2012/02/why-we-created-julia/>

Images: Twitter.com, Wikidata.org, Wikipedia.org

Why is Julia cool?



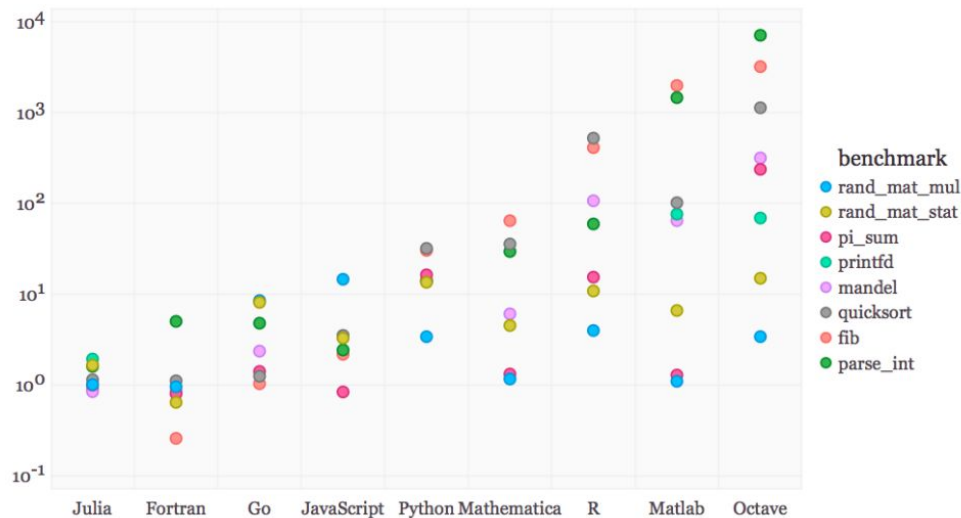
- Dynamically typed
- High-level syntax
- Built-in package manager
- Interactive development



- Great performance
- Works well with GPUs

Some cool features are...


- JIT compilation
 - Compilation happens at runtime
- Dynamic dispatch
 - Functional behavior depends on arguments
- Reflection and metaprogramming
 - Most of Julia is written in Julia
- Effective debugging
 - Debugger is a package
- Unicode integration
 - Gimmick?



Source: MIT.edu

So, you mentioned GPUs?

- Mature support for NVIDIA
 - CUDA.jl
- Similar but newer support for Intel
 - oneAPI.jl
- Experimental support for AMD
 - AMDGPU.jl
- **Kernels are Julia functions!**
 - ...they're not just translated into C code



Source: JuliaGPU.org

Choose your style... or both

Array Programming

```
julia> a = CuArray{Float32}(undef, (2,2));  
  
CURAND  
julia> rand!(a)  
2×2 CuArray{Float32,2}:  
0.73055  0.843176  
0.939997 0.61159  
  
CUBLAS  
julia> a * a  
2×2 CuArray{Float32,2}:  
1.32629 1.13166  
1.26161 1.16663  
  
CUSOLVER  
julia> LinearAlgebra.qr!(a)  
CuQR{Float32,CuArray{Float32,2}}  
with factors Q and R:  
Float32[-0.613648 -0.78958; -0.78958 0.613648]  
Float32[-1.1905 -1.00031; 0.0 -0.290454]
```

```
CUFFT  
julia> CUFFT.plan_fft(a) * a  
2-element CuArray{Complex{Float32},1}:  
-1.99196+0.0im  0.589576+0.0im  
-2.38968+0.0im -0.969958+0.0im  
  
CUDNN  
julia> softmax(real(ans))  
2×2 CuArray{Float32,2}:  
0.15712 0.32963  
0.84288 0.67037  
  
CUSPARSE  
julia> sparse(a)  
2×2 CuSparseMatrixCSR{Float32,Int32}  
with 4 stored entries:  
[1, 1] = -1.1905  
[2, 1] = 0.489313  
[1, 2] = -1.00031  
[2, 2] = -0.290454
```

Kernel Programming

Indexing

```
CUDA C: threadIdx.x; blockDim.y;  
CUDA.jl: threadIdx().x; blockDim().y
```

Cooperative groups

```
CUDA C: cudaLaunchCooperativeKernel(kernel, ...);  
CUDA.jl: @cuda cooperative=true kernel(...)
```

Shared memory

```
CUDA C: __shared__ int a[64];  
CUDA.jl: a = @cuStaticSharedMem{Int, 64}
```

```
CUDA C: extern __shared__ int b[];  
kernel<<<...,...,n*sizeof(int)>>>(...);  
CUDA.jl: b = @cuDynamicSharedMem{Int, 64}  
@cuda shmem=n*sizeof{Int} kernel(...)
```

Shuffle

```
CUDA C: __shfl_down_sync(mask, var, delta);  
CUDA.jl: shfl_down_sync(mask, var, delta)
```

Dynamic parallelism

```
CUDA C: kernel<...>(...);  
CUDA.jl: @cuda dynamic=true kernel(...)
```

Standard output

```
CUDA C: printf("Thread %d\n", threadIdx.x);  
CUDA.jl: @cuprintln("Thread ${threadIdx().x}")
```

Atomics

```
CUDA C: atomicAdd(ptr, val);  
CUDA.jl: @atomic a[...] += val
```

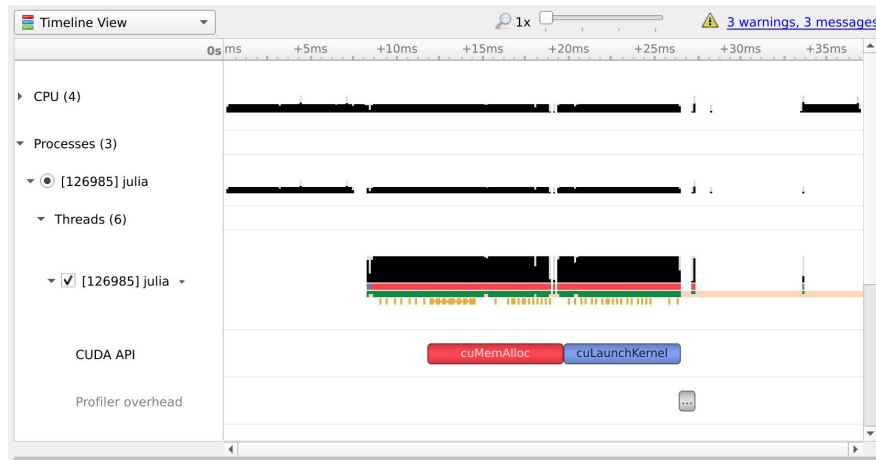
What about profiling and debugging?

Profiling

- Timing code is straightforward & robust
 - BenchmarkTools.jl
- nvprof
- nvvp
- Nsight tools

Debugging

- Code reflection macros
 - Interactively descend into call stack
 - Chthulhu.jl
- cuda-memcheck




Okay, so what's the catch?

- There's not really a catch...
 - I've had a positive experience working with Julia

However...

- When writing kernels...
 - Some Julia functionality is disabled
 - Dynamic typing, garbage collection, I/O
- Compilation + REPL start-up takes a while
 - Constant
- 1-based indexing is annoying
 - Personal preference



StickPNG.com



Code Demo

- The repo is linked in my project thread on Slack
 - https://code.vt.edu/neilg99/julia_exploration
- Live code?
- Questions?